

reworking your workflow

combining data analysis and reporting with dynamic documents

Christof Neumann

RTG Understanding Social Relationships

2021-11-18

content

- dynamic documents, huh?
- what is R markdown and why should we care?
- the dark side of a dynamic workflow
- let's get our hands dirty
- my assumptions:
 - you have used R before and use RStudio
 - you want to produce pdfs (article, thesis, presentation) or html (presentation, websites, shiny)
- adjust your expectations:
 - we scratch on the surface
 - always keep in mind what you **need** (and not what you **can do**)
 - fine line between what I would consider a 'good' presentation and showcasing the capabilities of R markdown

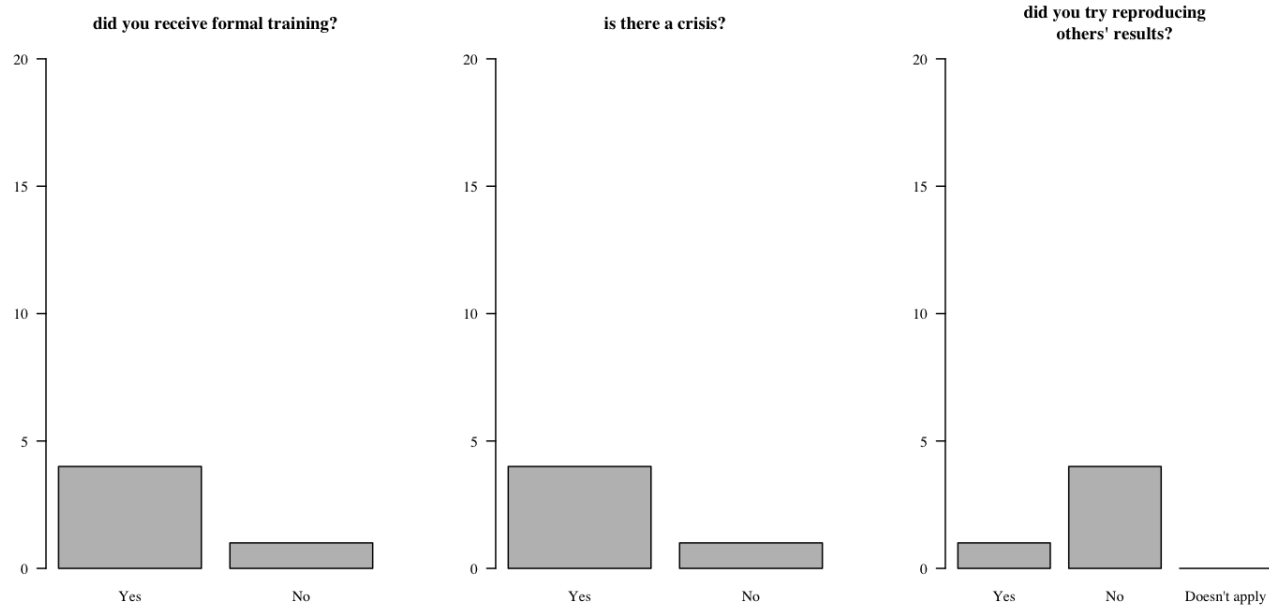
dynamic documents

dynamic (or living) documents

- continuous updates, revisions and edits (think Wikipedia)
- as opposed to static (dead) document (think printed lexicon)
- in scientific context:
 - preprints
 - during write-up
- can you do good and transparent science without it? YES.

a truly dynamic document

- this example grabs the data directly from the survey website



background

- to replicate

obtaining consistent results across studies aimed at answering the same scientific question, each of which has obtained its own data.

- to reproduce

obtaining consistent results using the same input data, computational steps, methods, and conditions of analysis

reproducibility


If someone applies the same analysis to the same data, then the same result should recur.

- process failure:
 - no results obtainable
 - unavailability of data and code
 - incomplete information about the analysis, software or tools
- outcome failure:
 - different results
 - error in original study or in reproduction attempt

what is R markdown and why should we
care?

R markdown?

- **markdown** is a programming language for converting plain text into formatted text
- **R markdown** builds on this and allows producing formatted output that builds on R code
- by now, the rmarkdown universe allows production of
 - text documents (journal articles, theses, books)
 - presentations
 - websites (via blogdown)
 - interactive applications/dashboards (via shiny)
 - exotic things (epub, docx, pptx, package documentation)




R markdown?

- all these have in common:
 - a shared logic and style for mixing written text and results from analyses
 - full reproducibility (at least in theory)
 - being open source and having an active community
 - still a high degree of flexibility
 - a certain aesthetic (it just looks good, no?)

R markdown!

- self-contained
- reproducible
- robust
- it's what the cool kids do these days (thanks, Roger)
- but there are downsides, too (we'll get to that)

why again?



I can't find the origin of this cartoon. If you know, please let me know.

**your most important collaborator is your
future self!!!**

when does it get complicated?

- it might very well...
- actually, it's likely to happen sooner rather than later
- long execution time (think bootstraps, models in `stan`)
 - sending workspaces to Roger
 - save workspace or not? it depends
- collaborators/supervisors/students won't/can't be bothered
- (in)dependence of RStudio
- environment [R version, package versions] at time of writing (when it works) versus time of production [which is the same as writing but also may be months or years later (reproducibility)]
- subtle differences (e.g. this presentation format doesn't support references and only very simple tables)

the nitty and the gritty

markdown (markup language)

- allows **plain text** with **simple formatting** to be translated into other formats
- independent of R, RStudio and *L^AT_EX*
- .md

Heading

=====

Sub-heading

Paragraphs ****are**** separated
by a blank line.

R markdown

- extension of markdown
- allows combining markdown ('text', 'narrative') and R ('code')
- still plain text
- independent of *L^AT_EX* and RStudio (although the latter provides lots of convenience)
- .Rmd

We can combine markdown and `R` code.
Backticks are really important here.

```
```${r, fig.cap = "My results."}  
plot(1:10)
```
```

knitr (and rmarkdown)

- are R packages
- knitr takes the code chunks, runs them and 'knits' them back into markdown
- requires R, but its results are independent of R (because they are markdown)
- rmarkdown and knitr also provide several convenience functions

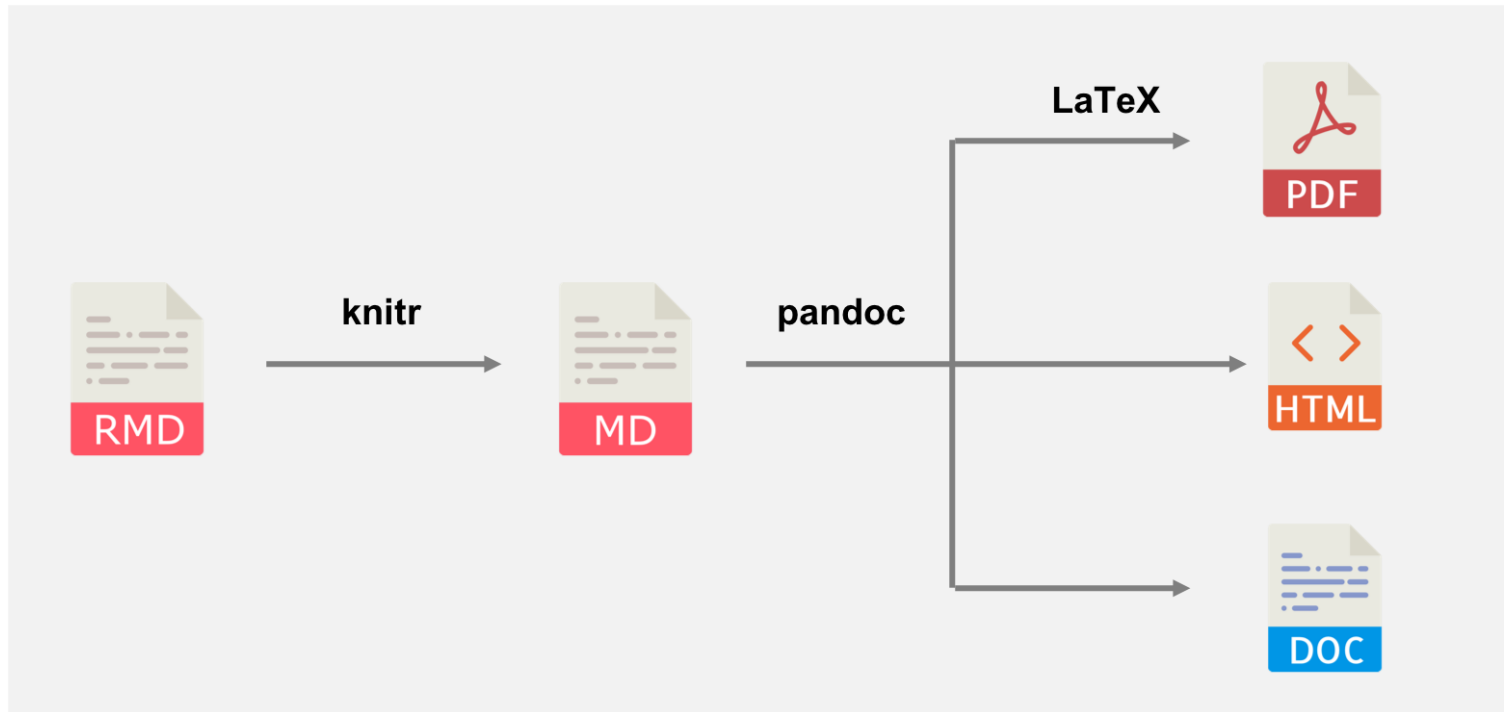
Pandoc

- document converter
- transforms .md into .html/.tex
- can also do tons of other conversions (-> MS Office, LibreOffice, ...)
- independent of R and RStudio

*L**A**T**E**X*

- builds .pdf files (article, thesis, presentation) from .tex files
- "a software system for document preparation"
- "separating presentation from content"
- independent of R and RStudio

in a nutshell



getting to it

prep

- create a new RStudio project
- create a new R markdown file
- memorize some keyboard shortcuts
 - cmd/ctrl + shift + k: build document
 - cmd/ctrl + option/alt + i: new R chunk
 - cmd/ctrl + shift + enter: run current chunk
 - cmd/ctrl + shift + 0: restart R

yaml header

- contains meta data for the final document (author, title, date)
- determines output format (here: `html_document`)
 - possibly with options (table of contents, highlighting options)
 - if we get to it, we might also try `pdf_document`
- other options:
 - control RStudio's behavior
 - bibliography data, reference style
 - line numbering and line spacing
- we won't look into this much, and mostly rely on the default options RStudio provides

yaml header

```
---  
title: "Untitled"  
author: "Christof Neumann"  
date: "9/29/2021"  
output:  
  html_document:  
    highlight: tango  
editor_options:  
  chunk_output_type: console  
---
```

a simple paragraph

```
# Introduction
```

A short paragraph that contains a summary of the article, for example the background and some major results. And note that the header 'Introduction' was formatted as a header because we put a \# in front of it (which is visible only in the source code of the document).

Introduction

A short paragraph that contains a summary of the article, for example the background and some major results. And note that the header 'Introduction' was formatted as a header because we put a # in front of it (which is visible only in the source code of the document).

simple formatting

Text can be ****bold**** or in **italics**.
But we can also use `~~strike through~~`.
Underlining is trickier and requires html code.

Text can be **bold** or in *italics*. But we can also use ~~strike through~~. Underlining is trickier and requires html code.

a first code chunk

This is just some text about how to generate random numbers in R:

```
```{r}  
rnorm(3)
```
```

Easy, no?

This is just some text about how to generate random numbers in R:

```
rnorm(3)
```

```
## [1] -0.61901958 -0.05424848 -0.04937414
```

Easy, no?

more on code chunks

```
```{r}  
x <- rnorm(1000)
hist(x)
```
```


```
x <- rnorm(1000)  
hist(x)
```

more on code chunks

- use chunk options

```
```{r, out.width = "40%", fig.align = 'center'}  
x <- rnorm(1000)
hist(x)
```
```


```
x <- rnorm(1000)  
hist(x)
```



more on code chunks

- hide the actual code, and add a caption (which is desired in many situations)

```
```{r, out.width = "35%", fig.align = 'center', echo = FALSE,  
fig.cap = "This is the figure caption."}
x <- rnorm(1000)
hist(x)
```
```



This is the figure caption.

using R objects in the actual text part

```
```{r, echo = FALSE}  
data("cars")
ncars <- nrow(cars)
mycor <- cor(cars$speed, cars$dist)
```
```

The cars data set has `r ncars` observations.

And the correlation between speed and distance is `r round(mycor, 2)`.

The cars data set has 50 observations. And the correlation between speed and distance is 0.81.

equations

We can also write fancy equations like $x = \sqrt{y} + c^2$, if they are wrapped in dollar signs (\$). But equations can also span multiple lines, with a 'proper' math environment:

```
$$  
x = \sum{\frac{42 + y}{\log(z)}}  
$$
```

Fancy, isn't it?

We can also write fancy equations like $x = \sqrt{y} + c^2$, if they are wrapped in dollar signs (\$). But equations can also span multiple lines, with a 'proper' math environment:

$$x = \sum \frac{42 + y}{\log(z)}$$

Fancy, isn't it?

tables

- are a bit trickier
- simple tables are possible with `knitr::kable()`
- more complex tables require the package `xtable`

```
```{r, echo = FALSE}
data("airquality")
knitr::kable(head(airquality))
```
```

| Ozone | Solar.R | Wind | Temp | Month | Day |
|-------|---------|------|------|-------|-----|
| 41 | 190 | 7.4 | 67 | 5 | 1 |
| 36 | 118 | 8.0 | 72 | 5 | 2 |
| 12 | 149 | 12.6 | 74 | 5 | 3 |
| 18 | 313 | 11.5 | 62 | 5 | 4 |
| NA | NA | 14.3 | 56 | 5 | 5 |
| 28 | NA | 14.9 | 66 | 5 | 6 |

references and bibliography

- (can't be demonstrated in this particular presentation format, bummer)
- requires `bibtex` references
- `bibtex` files are plain text files
- probably any reference manager can import/export references in `bib(la)tex` style (Zotero, Mendeley, Endnote)
- in my workflow, I manage my references with Zotero, and export regularly a `.bib` file with all the references I ever cited at some point, which I then use for citations in R markdown documents
- check out the "Better Bib(La)TeX" plugin for Zotero (<https://github.com/retorquere/zotero-better-bibtex>)
- a multi-platform graphic reference manager for `bib(la)tex` files is JabRef (<http://www.jabref.org/>)
- style sheets for different journals (`.csl` files) can be found here: <https://www.zotero.org/styles> (or here: <https://github.com/citation-style-language/styles>)

references and bibliography

- a typical bibtex entry for a journal article and book chapter (on the next slide) looks like this
- the crucial thing here is the citation key, i.e. the identifier (forstmeier2011 in the first example)
- you need the citation key to cross-reference

```
@article{forstmeier2011,  
  title = {Cryptic Multiple Hypotheses Testing in Linear Models:  
    Overestimated Effect Sizes and the Winner's Curse},  
  volume = {65},  
  doi = {10.1007/s00265-010-1038-5},  
  journaltitle = {Behavioral Ecology and Sociobiology},  
  date = {2011},  
  pages = {47--55},  
  author = {Forstmeier, Wolfgang and Schielzeth, Holger}  
}
```

references and bibliography

- here is an entry for a book chapter

```
@incollection{garamszegi2014,  
  location = {Berlin},  
  title = {Multimodel-Inference in Comparative Analyses},  
  booktitle = {Modern Phylogenetic Comparative Methods and  
    Their Application in Evolutionary Biology},  
  publisher = {Springer},  
  date = {2014},  
  pages = {305--331},  
  author = {Garamszegi, László Zsolt and Mundry, Roger},  
  editor = {Garamszegi, László Zsolt}  
}
```

references and bibliography

- update the yaml header to point to .bib file (biblographic entries) and .csl file (formatting style)

```
---
title: "Untitled"
author: "Christof Neumann"
date: "9/29/2021"
output:
  html_document:
    highlight: tango
editor_options:
  chunk_output_type: console
bibliography: data_files/refs.bib
csl: data_files/animal-behaviour.csl
---
```

references and bibliography

```
Blah blah [@forstmeier2011; @garamszegi2014].
```

```
Blah blah [see @forstmeier2011, p. 365; but see @garamszegi2014].
```

```
@forstmeier2011 say something.
```

- bibliography is inserted after the header # references

wrapping up

going further

- cross-references (to figures and tables)
- more table and plotting options
- conditional inclusion of chunks
- actually using *L^AT_EX*
- caching of code chunks
- document formatting/layouts (e.g. thesis chapters)
- extended referencing (bibliographic)
- integrating other coding languages (Python, stan, C++, ...)
- put it all under version control

wrapping up

- establish good practice
- you can't reproduce what doesn't exist
- think about your workflow
- your single most important collaborator is your future self